

CS 437: Internet of Things

Final Project Report

Team

- Eric Schrock (ejs9)
- Devin Schmitz (devinmms3)

Overview

We created a distributed network of CO₂ sensors that report back via wifi to a central hub. The hub exposes a web interface on the LAN for accessing the data. This allows us to track variations in CO₂ levels throughout a home and over time. This data could be used to understand and improve indoor air quality, leading to healthier lives.

Motivation

Breathing is one of the most fundamental processes of our lives, but most homes are only equipped with smoke and carbon monoxide detectors, which just warn against critically dangerous levels of air pollution. High levels of pollutants like Carbon Dioxide (CO₂), particulate matter (PM), and volatile organic compounds (VOCs) can go unnoticed, contributing to drowsiness, illness, and poor cognitive function, among many other potential poor health outcomes¹²³.

Large pollutants can be filtered out by a good HVAC or air filtration system, but small asphyxiants such as CO₂ and VOCs are best managed with proper air ventilation. However, the exchange of fresh outdoor air for indoor air works against the mission of modern HVAC, door, and window systems, which attempt to maintain a consistent internal climate (in spite of the surrounding environment) by *avoiding* the exchange of indoor and outdoor air.

It is easy to open a window during the spring and fall when the weather is nice, but knowing when it is necessary to open a window for proper air circulation during the cold winter or hot summer can be a challenge. Having a window open

¹<https://ehp.niehs.nih.gov/doi/10.1289/ehp.1510037>

²https://papers.ssrn.com/sol3/papers.cfm?abstract_id=4296077

³<https://www.niehs.nih.gov/health/topics/agents/air-pollution/index.cfm>

all the time wastes energy heating excess fresh air, but fatigue and headache from poor air quality only exacerbates common ailments like seasonal depression.

The first step to balancing this tension is having data. CO₂ serves as a good proxy for appropriate air ventilation in general. CO₂ is a direct byproduct of respiration and combustion, so indoor levels rise throughout the day as we breath and use gas appliances like stoves, hot water heaters, and furnaces. By monitoring indoor CO₂ levels, we can have a better idea of when there is adequate ventilation and when more fresh air exchange is needed.

However, proper air monitoring requires a distributed network. A single sensor, though useful, does not provide a comprehensive picture of the air quality throughout a home (especially in homes with poor overall ventilation). A network of sensors also reduces inaccuracy due to any single faulty sensor reading and keeps the system more robust to sensor failures. Additionally, many existing air monitoring solutions are cost prohibitive to deploy across a home, with individual sensors costing up to \$230⁴ and coming with a list of extra and unnecessary features like onboard data processing and display.

We designed a simple, low cost CO₂ sensor that can be deployed throughout a home for real time monitoring of air quality with a single central hub for processing and displaying the data via a convenient web application on the local area network.

Technical Approach

Our overall system consists of a base station connected to a network of sensors over wifi. The base station is a Raspberry Pi 4B running a server. Sensors connect to the server and periodically report CO₂ measurements. The base station then presents those measurements on a website available on the local area network.

Each sensor consists of a Raspberry Pi Pico WH connected to a CO₂ sensor over I₂C. Each sensor runs a client that is responsible for connecting to the server on the base station, reading CO₂ measurements, and sending those measurements to the base station.

⁴<https://www.airthings.com/wave-plus>

⁵The CO₂ sensor shown in figure 2 does not match the actual part and is only meant to give a general sense of the design.

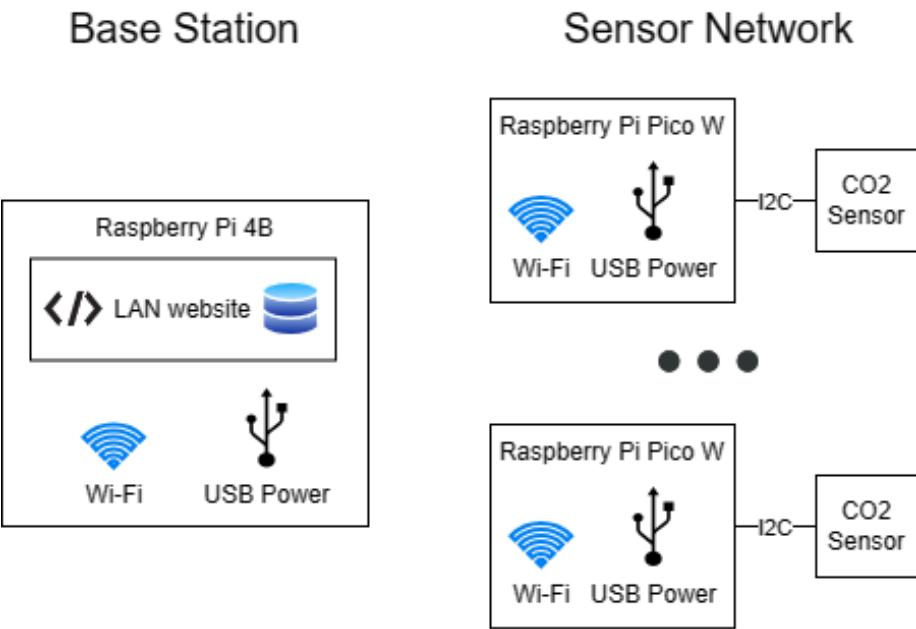


Figure 1: System Architecture

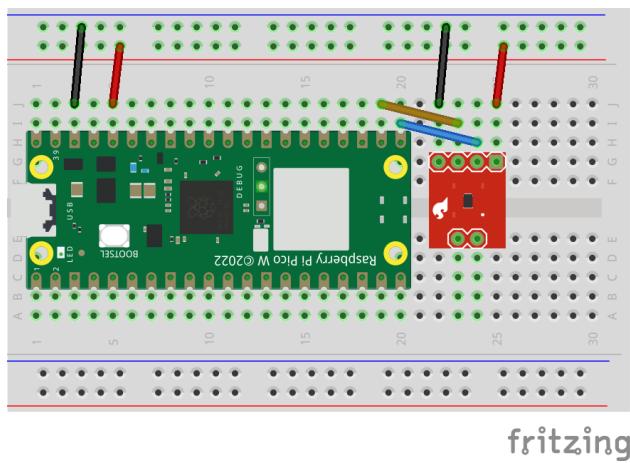


Figure 2: Sensor Layout⁵

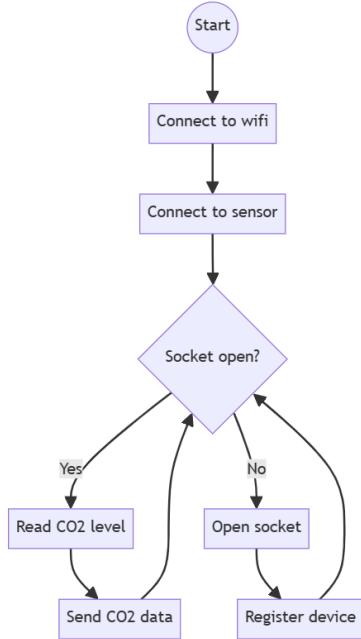


Figure 3: Sensor Flowchart

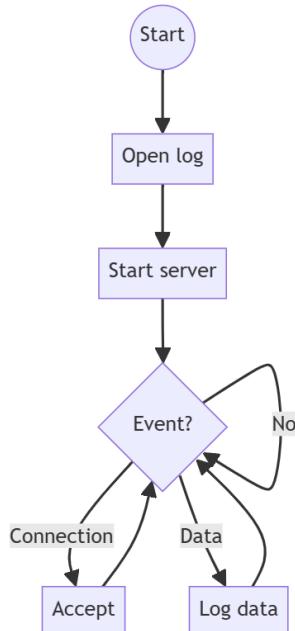


Figure 4: Server Flowchart



Figure 5: Website Host Flowchart

Implementation Details

Sensors

We spent a significant amount of time comparing⁶ different options for the microcontroller on our sensors units. We started our search with six different Arduino variants (five Nano variants and the Micro). We looked at price, availability, breadboard compatibility, power and IO options, and connectivity (wifi, Bluetooth, and BLE). We were leaning towards the Arduino Nano 33 IoT⁷ when we found the Raspberry Pi Pico WH^{8 9}, which checked all our boxes but for a much cheaper price (\$7 vs \$27).

The Pico comes with both C/C++¹⁰ and Python¹¹ SDKs. We are both embedded C programmers for our day job, so we picked the Python SDK to get some experience with MicroPython¹². Additionally, the simplicity of Python fit with a simple prototype on a tight timeline.

⁶<https://github.com/EricSchrock/co2-monitor/blob/main/docs/microcontroller.md>

⁷<https://store-usa.arduino.cc/products/arduino-nano-33-iot-with-headers>

⁸<https://www.raspberrypi.com/documentation/microcontrollers/raspberry-pi-pico.html>

⁹<https://www.pishop.us/product/raspberry-pi-pico-wh-pre-soldered-headers>

¹⁰<https://datasheets.raspberrypi.com/pico/raspberry-pi-pico-c-sdk.pdf>

¹¹<https://datasheets.raspberrypi.com/pico/raspberry-pi-pico-python-sdk.pdf>

¹²<https://micropython.org/>

We also considered¹³ multiple CO₂ sensors, before settling on the ENS160¹⁴. Our criteria were price, availability, breadboard compatibility, and power and IO options.

Additionally, we had to consider whether to choose a true CO₂ sensor or an equivalent CO₂ (eCO₂)¹⁵ sensor. eCO₂ sensors don't measure CO₂ directly. Instead, they measure total volatile organic components (TVOC) and use that measurement to estimate the CO₂ level. eCO₂ sensors are much cheaper but are less accurate. We chose to go with an eCO₂ sensor, as price is a big factor in our design. Having multiple sensor units allowed us to sanity check the reported values and notice anomalies. Given this, the eCO₂ sensor was accurate enough to observe trends and draw conclusions. More on this in the CO₂ Data Analysis portion of the Results section.

The SparkFun ENS160¹⁶ came with the option to communicate over I2C or SPI. We chose I2C because it was simpler to implement in MicroPython.

The Pico supports wifi, Bluetooth, and BLE. We initially planned to use BLE to communicate with the server, but implementing wifi communication turned out to be much simpler. Additionally, power usage turned out to be less of a concern than we initially thought because we chose to power the sensor units off wall power instead of batteries.

We chose to connect up the Pico and ENS160 on a half sized breadboard and to power them both through the micro USB port on the Pico. We built both dev units and full prototypes.

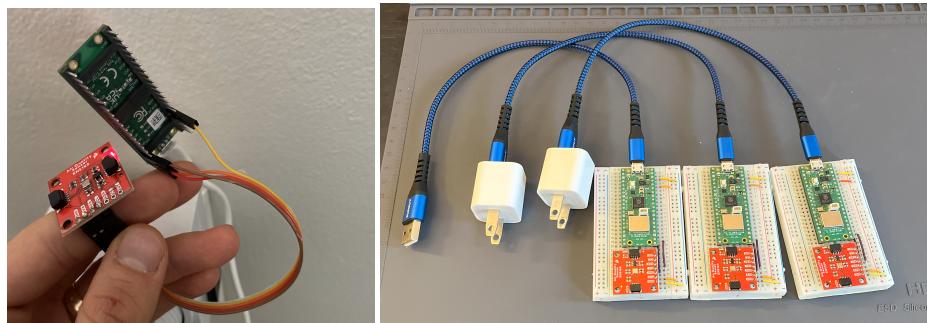


Figure 6: Development Unit (left) and Full Prototypes (right)

The software development for the sensor unit was relatively straight forward. We wrote simple programs to aid hardware bring up, first to blink the Pico LED

¹³<https://github.com/EricSchrock/co2-monitor/blob/main/docs/co2-sensor.md>

¹⁴<https://www.sciosense.com/wp-content/uploads/documents/SC-001224-DS-9-ENS160-Datasheet.pdf>

¹⁵<https://electronics360.globalspec.com/article/17986/what-are-eco2-sensors>

¹⁶<https://www.sparkfun.com/products/20844>

and then to turn the Pico LED on when we breathed on the CO₂ sensor. Then we created a simple program that connects to wifi, connects to the CO₂ sensor via I₂c, connects to the server, and then periodically reads and reports the CO₂ level (see the sensor flow chart above in the Technical Approach section).

Server

In choosing the hardware for the central server, we quickly decided to use the Raspberry Pi 4B¹⁷ that we had both already purchased for earlier labs in the class. The Raspberry Pi 4B provides a builtin wifi module to connect to the sensor units and sufficient processing power and RAM to host a basic webpage and do minimal data processing. In fact, the Raspberry Pi 4B far exceeded our basic requirements for a server. If we were going to productionize the system, we would explore using a smaller Raspberry Pi with less RAM or potentially an even simpler SBC to reduce the cost of the system.



Figure 7: Raspberry Pi Hubs Connected to Sensors via Wifi

In designing the server software, our goal was simplicity, reliability, and minimal dependencies. We wrote the entire server application using just the Python standard library. We started with a simple server that could connect to one sensor using sockets, using what we had learned from our previous labs. Then, we extended the server and our knowledge of sockets programming by incorporating a selector¹⁸ to handle multiple socket connections to different sensors. Throughout the server development, we took advantage of the ability to create socket pairs on a single device to emulate a sensor network locally for debugging, which allowed us to parallelize the sensor and server development.

When it came to storing CO₂ data, we considered integrating a proper database.

¹⁷<https://www.raspberrypi.com/products/raspberry-pi-4-model-b/>

¹⁸<https://docs.python.org/3/library/selectors.html>

However, after looking into it, even a simple SQL database¹⁹ seemed like unnecessary overhead and complexity for the data we were storing. We settled on storing the data in basic CSV files to make both writing the data and loading it for later analysis easy. Our first iteration of data logging had a unique file for each sensor, which contained the full date-time timestamp for each CO2 reading. It quickly became apparent that this was suboptimal for two reasons. First, having all CO2 data from one sensor in a single file meant that significantly more data would have to be parsed to analyze a small time window of data. Second, including the date in every timestamp was a significant duplication of data and thus waste of memory. We settled on splitting up the data by date instead, so that each day's readings would be stored in a directory with a CSV file for each sensor that only contained the timestamps for the CO2 readings.

One challenge I (Eric) faced was the server process on the Pi dying every few hours. My solution was to reboot the Pi every hour via the Linux `cron` utility. I edited `/etc/rc.local` to restart the server on each reboot, allowing me to capture data over long periods without babysitting (e.g overnight). Devin's system never exhibited this issue.

Website

Since neither of us has web development experience, we were learning everything as we went. This meant simplicity and basic functionality were high priorities as we chose software stacks for our web application. We are both comfortable with Python, so it made sense to use a backend Python framework. We considered two dominant web frameworks: Django²⁰ and Flask²¹. Each has pros and cons, and both are widely used in industry. However, we settled on Flask for two key reasons. First, it is faster to get a minimal web application up and running from scratch in Flask. Second, Flask uses `jinja2`²² templates by default, which we were already familiar with.

From there, the development of the website evolved in stages. First, we setup a basic “Hello World” page to verify that we could successfully run our Flask application. Then, we upgraded the application to include a request form for a date, which we initially just echoed to the webpage. Finally, we updated the web app to display the list of CO2 reading times and values for the selected date.

At this stage, we had to decide on a software stack to generate the most important part of the web application: a chart of the data. We considered three

¹⁹<https://docs.python.org/3/library/sqlite3.html>

²⁰<https://docs.djangoproject.com/en/5.0/>

²¹<https://flask.palletsprojects.com/en/3.0.x/>

²²<https://jinja.palletsprojects.com/en/3.1.x/templates/>

options: Chart.js²³, Matplotlib²⁴, and plotlib²⁵ with WebAssembly²⁶. These three options vary significantly, from the language used (Javascript, Python, and Rust respectively) to how they integrate into the Flask application. While the Rust and WebAssembly option was interesting to explore, WebAssembly is a relatively new technology and doesn't have the same level of existing examples and documentation as the more established competitors. Chart.js was a compelling option, with support for dynamic and interactive charts in a web-native language. However, Matplotlib had a few great examples for integrating charts into a webpage and was already familiar to us, so it won out as the best option to meet the goals and timelines of this project. Additionally, in brief prototyping, a webpage with a matplotlib image reloaded nearly instantaneously, while plotting the same data with Chart.js had a noticeable (several second) delay.

Finally, we needed to choose a WSGI server to host our web application. We considered several of the options on Flask's list of recommended self-hosted WSGI servers²⁷. We chose Waitress²⁸ because it is a pure Python application with no dependencies, which made it easy to include in our project.

Results

Project Timeline

Overall, the project probably took a little over the projected 45 hours. The bookends, the initial design decisions and the project report, took much more time than planned, but the hardware bring up was exceptionally smooth and we did not need the planned second prototype. Our time estimates for the website were about right.

Project Objectives

At a high level, we met the objective we laid out in our project proposal, which was “to create a network of CO₂ sensors with a web interface” to allow “distributed monitoring of indoor air quality.” Our proposal also laid out the objectives of “simple, low cost, low energy” sensors.

We met the objective of simple sensors. We avoided extra functionality that would have increased cost and complexity (e.g. multi-function sensors or extra website features). We also kept the setup instructions²⁹ simple (for a prototype).

As our design evolved from a battery powered sensor to a wall powered sensor, we dropped the objective of being low energy. This allowed us to make tradeoffs in favor of simplicity and speed. For example, it allowed us to keep our code

²³<https://www.chartjs.org/>

²⁴<https://matplotlib.org/stable/index.html>

²⁵<https://docs.rs/plotlib/latest/plotlib/>

²⁶<https://rustwasm.github.io/>

²⁷<https://flask.palletsprojects.com/en/3.0.x/deploying/>

²⁸<https://docs.pylonsproject.org/projects/waitress/en/stable/>

²⁹<https://github.com/EricSchrock/co2-monitor/blob/main/src/README.md>

simple by using wifi instead of BLE and by not implementing any low power modes. Now that we have a working prototype that has proved it's value, power usage optimization would be a valuable consideration for a second prototype.

We also met the objective of low cost sensors, with each sensor unit costing roughly \$40. A productionized sensor unit, benefiting from economies of scale, could cost even less.

The base station prototype is pricier at \$88. This is largely because we used the 8 GB Raspberry Pi 4B because it is what we had on hand. Using the 1 GB model would drop the base station cost to \$48. A simple website doesn't take much compute power. A productionized version of the base station could use an even cheaper processor. It's even possible that one of the sensor units could function as the server for the rest of the sensor network.

Sensor Unit

Part	Price ³⁰
Raspberry Pi Pico WH	\$7.00 ³¹
12.5W power supply	\$8.00 ³²
Sparkfun ENS160	\$19.95 ³³
Half sized bread board	\$4.75 ³⁴
Wires and headers	~\$0.30
Total	\$40.00

Base Station

Part	Price ³⁵
Raspberry Pi 4B (8 GB)	\$75.00 ³⁶
15W power supply	\$8.00 ³⁷
Case	\$5.00 ³⁸
Total	\$88.00

³⁰Price does not include shipping or bulk discounts.

³¹<https://www.pishop.us/product/raspberry-pi-pico-wh/>

³²<https://www.pishop.us/product/raspberry-pi-12-5w-power-supply-us-white/>

³³<https://www.sparkfun.com/products/20844>

³⁴<https://www.pishop.us/product/half-size-400-pin-diy-breadboard-white/>

³⁵Price does not include shipping or bulk discounts.

³⁶<https://www.pishop.us/product/raspberry-pi-4-model-b-8gb/>

³⁷<https://www.pishop.us/product/raspberry-pi-15w-power-supply-us-white/>

³⁸<https://www.pishop.us/product/raspberry-pi-4-case-red-white/>

CO2 Data Analysis

One question we had was whether our sensors would be accurate enough to provide useful data. We found that the data trends were reliable enough to interpret and act on. For example, figure 8 shows around thirteen hours of CO2 data that is readily attributable to the day's activities.

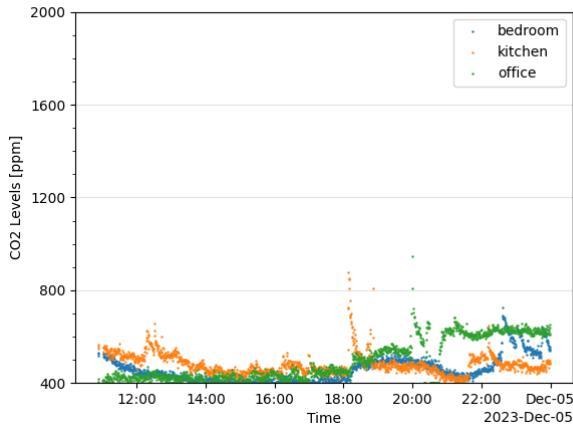


Figure 8: CO2 Data from Eric's Home (12/4/23)

1. 11 AM - 2 PM: Bedroom CO2 level drops from its nighttime high.
2. 12 - 1 PM: Kitchen CO2 level jumps as my wife and I (Eric) eat lunch.
3. 1 - 6:15 PM: Bedroom sits empty and has the lowest CO2 level. I sit at my desk in my office and my wife moves around the kitchen and adjoining living room. The kitchen CO2 level is higher than the office CO2 level because humans breath out more CO2 when we are active.
4. 6:15 PM: Kitchen CO2 level spikes when we vent the pressure cooker and then dissipate to the rest of the house.
5. 6:30 - 7 PM: Kitchen CO2 level jumps as my wife and I eat supper.
6. 7 - 9:30 PM: My wife leaves for time with friends and I return to my office. Kitchen CO2 level drops below that of the office.
7. 8:15 PM: I briefly open a window in my office. The office CO2 level drops to the minimum sensor value of 400 ppm, but jumps back up as soon as I close the window.
8. 8:30 - 8:45 PM: I open a window in my office for 15 minutes. The office CO2 level again drops to the minimum, but jumps back as soon as the window is closed.
9. 9:30 - 10:15 PM: My wife returns and we spend time together in the living room, near the kitchen. The kitchen CO2 level jumps.
10. 10:15 PM: We go to bed. Kitchen CO2 level drops and bedroom CO2 level jumps.

The sensors do have a couple issues to keep in mind. First, as explained in the section on implementation details, they are measuring eCO₂ which is an estimate of CO₂ levels derived from TVOC levels. TVOCs can spike without CO₂ spiking. For example, the data above shows a spike in the office CO₂ levels at around 8 PM. This was from running a paper shredder near the sensor.

Second, the sensor results can occasionally get stuck at a high offset until they are power cycled. The CO₂ level in the office jumped significantly the evening of December 4th (figure 8) and remained high into the next day (figure 9). I first tried opening windows throughout the house from 10:15 - 10:30 AM. The other rooms responded to this, but the office CO₂ level jumped back up as soon as the windows were closed. At around noon, I power cycled the office sensor.

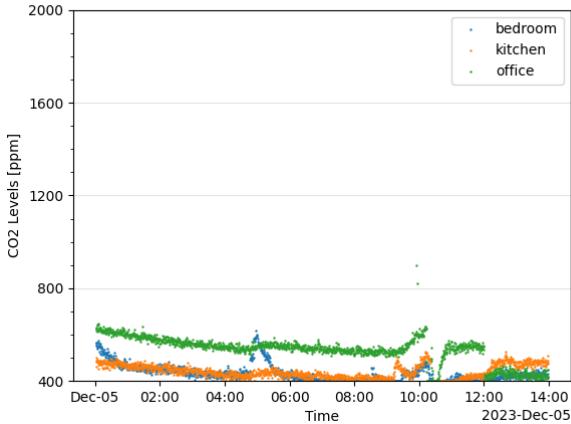


Figure 9: CO₂ Data from Eric's Home (12/5/23)

There were other interesting trends we observed in the data. For example, the week of Thanksgiving I (Devin) was out of the home, and all three CO₂ sensors consistently measured low and unchanging CO₂ levels (figure 10). This indicates that practically all significant CO₂ production in the home is linked to my activity in some way, rather than to furnace or hot water heater cycles.

Another pattern I noticed is a cyclic rising and falling CO₂ levels in the bedroom during early morning hours (figure 11). The cycles appear every night and only in the bedroom, leading me to wonder if they are related to sleep cycles. Correlating the data to a sleep tracker would be a fascinating follow on project.

Despite some issues, our sensor data correlated well to real world events. As long as the limitations we listed are kept in mind, we believe it is accurate enough to drive meaningful decisions. Examples range from whether to open a window in the winter to whether to invest in a new HVAC that cycles in outdoor air in response to high indoor CO₂ levels.

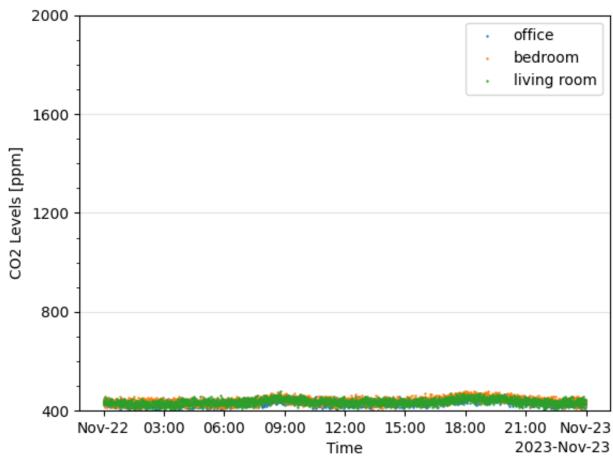


Figure 10: CO2 from Devin's Vacant Home (11/22/23)

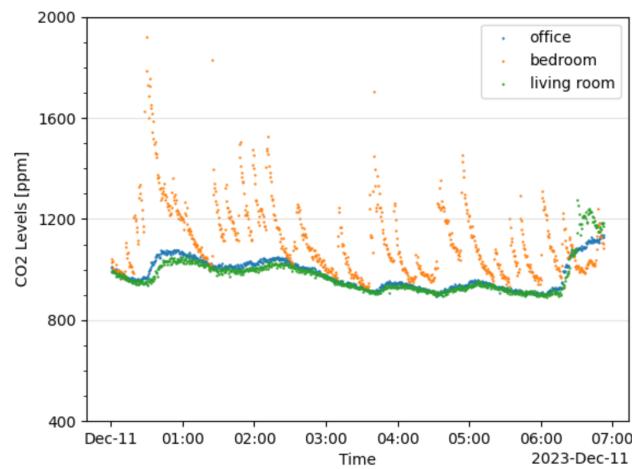


Figure 11: Cycling CO2 from Devin's Early Morning Bedroom

Demo Videos

We recommend watching these videos on a large screen and/or setting your video player to HD/1080p.

- [Hardware Demo](#)
- [Server Demo](#)
- [Website Demo](#)

Project Repository

<https://github.com/EricSchrock/co2-monitor>

What We Learned

Eric

Two of the key skills I learned in this project were networking via the Python sockets library and how to configure Linux using the `/etc/rc.local` file and cron jobs. Both were completely new to me and are valuable experiences for the future.

On a different note, I was surprised by how easy it was to get the project up and running on the Raspberry Pi Pico with MicroPython. I'm used to embedded boards being much more painful to bring up. Perhaps some personal projects I thought would be too time consuming are actually in reach!

This project also taught me about the health impacts of high concentrations of CO₂ in indoor spaces. This is important knowledge to help me stay sharp. As a remote worker, I can put what I've learned to use by opening my office windows at key times.

Devin

I learned a great deal from working with the hardware in this lab. I soldered header pins for the first time. I gained valuable experience reading the ENS160 datasheet to understand how to read sensor values over I₂C. I learned how to use `minicom` to view serial data from the Pi Pico and how to prototype micropython on the Pi Pico from the interactive micropython REPL.

This project significantly expanded my understanding and confidence in socket programming, which I had only just learned in previous class labs. It also challenged me to write a safe and robust server application.

Building the webpage may have taught me the most. Weighing different Python web frameworks gave great insight into the infrastructure underlying websites and web applications that I use on a daily basis. Actually deploying a functional web form involved more pieces than I knew and gave me a new respect for web development.

If We Had More Time

If we were to take this project further, we would have three main goals. The first would be to increase the accuracy of the sensor readings, either by tuning them with temperature and humidity readings reported to them by the base station or by automatically power cycling the ENS160 CO₂ sensor periodically.

Second, we would expand the functionality of the website. We would add the ability to save timestamped notes to mark and explain phenomena in the data. We would also make the time interval for the display configurable. Additionally, we would add simple statistics, such as the max and average over the selected time interval.

The third goal would be to condense the sensor unit into a wall wart with a protective housing. We would look into 3D printing for the case and into a custom PCB to fit in a smaller form factor.

Conclusion

In conclusion, we accomplished the high level goal of our project, to prototype a distributed, networked, and affordable CO₂ monitoring solution. Along the way, we tried new technologies, built new skills, and gained a deeper understanding of the quality of the air in our homes and of how it impacts our health.

Todo

todo (Eric): Record a demo video of how to setup the hardware.

- Sensor and base station hardware
- Installation (<https://github.com/EricSchrock/co2-monitor/blob/main/src/README.md>)
- Talk through LED transitions
- Show server log and CO₂ data files