

# CS 437: Internet of Things

## Final Project Report

### Team

- Eric Schrock (ejs9)
- Devin Schmitz (devinmms3)

### Overview

We created a distributed network of CO<sub>2</sub> sensors that report back via wifi to a central hub. The hub exposes a web interface on the LAN for accessing the data. This allows us to track variations in CO<sub>2</sub> levels throughout a home and over time. This data could be used to understand and improve indoor air quality, leading to healthier lives.

### Motivation

todo: (Eric) Review this. Does it need to change at all? It is copy-paste from the proposal, but I slightly modified the final paragraph.

Breathing is one of the most fundamental processes of our lives, but most homes are only equipped with smoke and carbon monoxide detectors, which just warn against critically dangerous levels of air pollution. High levels of pollutants like Carbon Dioxide (CO<sub>2</sub>), particulate matter (PM), and volatile organic compounds (VOCs) can go unnoticed, contributing to drowsiness, illness, and poor cognitive function, among many other potential poor health outcomes<sup>123</sup>.

Large pollutants can be filtered out by a good HVAC or air filtration system, but small asphyxiants such as CO<sub>2</sub> and VOCs are best managed with proper air ventilation. However, the exchange of fresh outdoor air for indoor air works against the mission of modern HVAC, door, and window systems, which attempt to maintain a consistent internal climate (in spite of the surrounding environment) by *avoiding* the exchange of indoor and outdoor air.

---

<sup>1</sup><https://ehp.niehs.nih.gov/doi/10.1289/ehp.1510037>

<sup>2</sup>[https://papers.ssrn.com/sol3/papers.cfm?abstract\\_id=4296077](https://papers.ssrn.com/sol3/papers.cfm?abstract_id=4296077)

<sup>3</sup><https://www.niehs.nih.gov/health/topics/agents/air-pollution/index.cfm>

It is easy to open a window during the spring and fall when the weather is nice, but knowing when it is necessary to open a window for proper air circulation during the cold winter or hot summer can be a challenge. Having a window open all the time wastes energy heating excess fresh air, but fatigue and headache from poor air quality only exacerbates common ailments like seasonal depression.

The first step to balancing this tension is having data. CO<sub>2</sub> serves as a good proxy for appropriate air ventilation in general. CO<sub>2</sub> is a direct byproduct of respiration and combustion, so indoor levels rise throughout the day as we breath and use gas appliances like stoves, hot water heaters, and furnaces. By monitoring indoor CO<sub>2</sub> levels, we can have a better idea of when there is adequate ventilation and when more fresh air exchange is needed.

However, proper air monitoring requires a distributed network. A single sensor, though useful, does not provide a comprehensive picture of the air quality throughout a home (especially in homes with poor overall ventilation). A network of sensors also reduces inaccuracy due to any single faulty sensor reading and keeps the system more robust to sensor failures. Additionally, many existing air monitoring solutions are cost prohibitive to deploy across a home, with individual sensors costing up to \$230<sup>4</sup> and coming with a list of extra and unnecessary features like onboard data processing and display.

We designed a simple, low cost CO<sub>2</sub> sensor that can be deployed throughout a home for real time monitoring of air quality with a single central hub for processing and displaying the data via a convenient web application on the local area network.

## Technical Approach

Our overall system consists of a base station connected to a network of sensors over wifi. The base station is a Raspberry Pi 4B running a server. Sensors connect to the server and periodically report CO<sub>2</sub> measurements. The base station then presents those measurements on a website available on the local area network.

Each sensor consists of a Raspberry Pi Pico WH connected to a CO<sub>2</sub> sensor over I2C. Each sensor runs a client that is responsible for connecting to the server on the base station, reading CO<sub>2</sub> measurements, and sending those measurements to the base station.

Note: The CO<sub>2</sub> sensor shown in the Fritzing diagram above does not match the actual part and is only meant to give a general sense of the design.

todo: Add a high-level diagram of the server code (details belong in the next section)

todo: Add a high-level diagram of the website code (details belong in the next section)

---

<sup>4</sup><https://www.airthings.com/wave-plus>

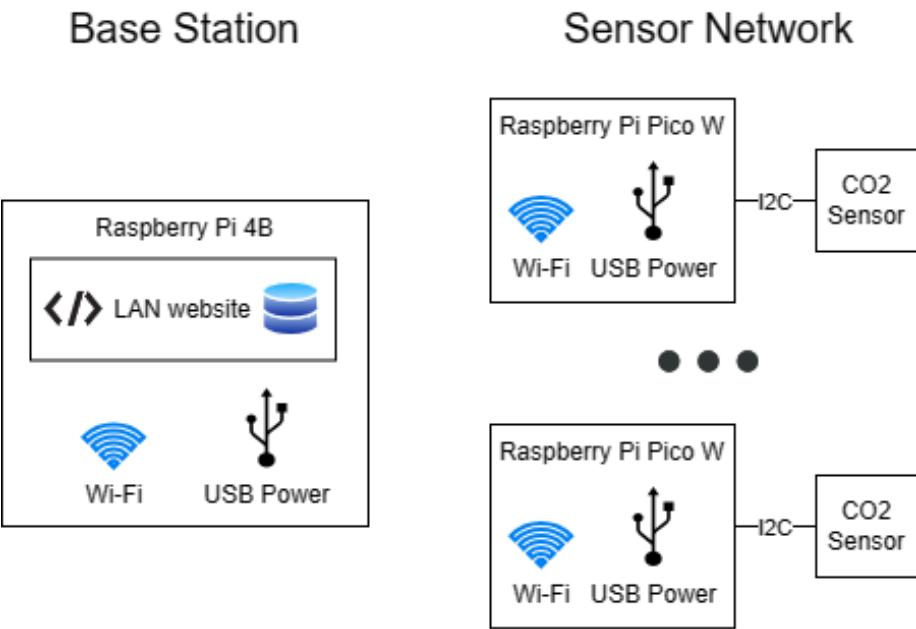


Figure 1: System Architecture

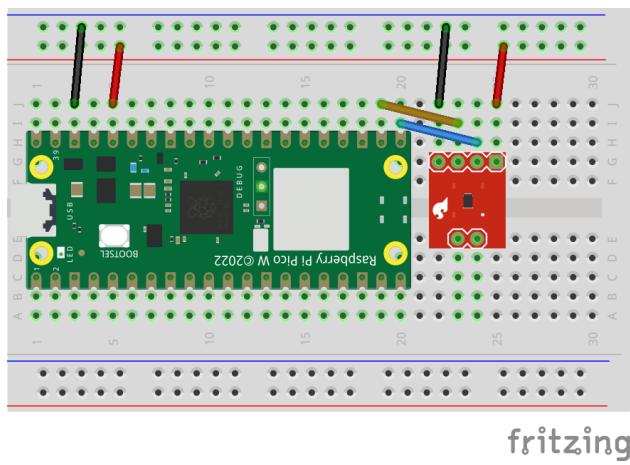


Figure 2: Sensor Layout

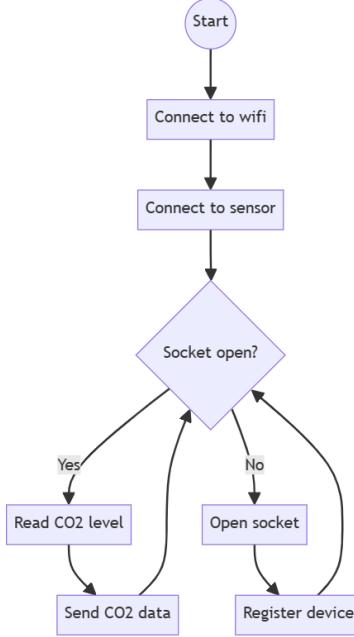


Figure 3: Sensor Flowchart

## Implementation Details

### Sensors

We spent a significant amount of time comparing<sup>5</sup> different options for the microcontroller on our sensors units. We started our search with six different Arduino variants (five Nano variants and the Micro). We looked at price, availability, breadboard compatibility, power and IO options, and connectivity (wifi, Bluetooth, and BLE). We were leaning towards the Arduino Nano 33 IoT<sup>6</sup> when we found the Raspberry Pi Pico WH<sup>7</sup> <sup>8</sup>, which checked all our boxes but for a much cheaper price (\$7 vs \$27).

The Pico comes with both C/C++<sup>9</sup> and Python<sup>10</sup> SDKs. We are both embedded C programmers for our day job, so we picked the Python SDK to get some experience with MicroPython<sup>11</sup>. Additionally, the simplicity of Python fit with a simple prototype on a tight timeline.

<sup>5</sup><https://github.com/EricSchrock/co2-monitor/blob/main/docs/microcontroller.md>

<sup>6</sup><https://store-usa.arduino.cc/products/arduino-nano-33-iot-with-headers>

<sup>7</sup><https://www.raspberrypi.com/documentation/microcontrollers/raspberry-pi-pico.html>

<sup>8</sup><https://www.pishop.us/product/raspberry-pi-pico-wh-pre-soldered-headers>

<sup>9</sup><https://datasheets.raspberrypi.com/pico/raspberry-pi-pico-c-sdk.pdf>

<sup>10</sup><https://datasheets.raspberrypi.com/pico/raspberry-pi-pico-python-sdk.pdf>

<sup>11</sup><https://micropython.org/>

We also considered<sup>12</sup> multiple CO<sub>2</sub> sensors, before settling on the ENS160<sup>13</sup>. Our criteria were price, availability, breadboard compatibility, and power and IO options.

Additionally, we had to consider whether to choose a true CO<sub>2</sub> sensor or an equivalent CO<sub>2</sub> (eCO<sub>2</sub>)<sup>14</sup> sensor. eCO<sub>2</sub> sensors don't measure CO<sub>2</sub> directly. Instead, they measure total volatile organic components (TVOC) and use that measurement to estimate the CO<sub>2</sub> level. eCO<sub>2</sub> sensors are much cheaper but are less accurate. We chose to go with an eCO<sub>2</sub> sensor, as price is a big factor in our design. Having multiple sensor units allowed us to sanity check the reported values and notice anomalies. Given this, the eCO<sub>2</sub> sensor was accurate enough to observe trends and draw conclusions. More on this in the CO<sub>2</sub> Data Analysis portion of the Results section.

The SparkFun ENS160<sup>15</sup> came with the option to communicate over I2C or SPI. We chose I2C because it was simpler to implement in MicroPython.

The Pico supports wifi, Bluetooth, and BLE. We initially planned to use BLE to communicate with the server, but implementing wifi communication turned out to be much simpler. Additionally, power usage turned out to be less of a concern than we initially thought because we chose to power the sensor units off wall power instead of batteries.

We chose to connect up the Pico and ENS160 on a half sized breadboard and to power them both through the micro USB port on the Pico. We built both dev units and full prototypes.

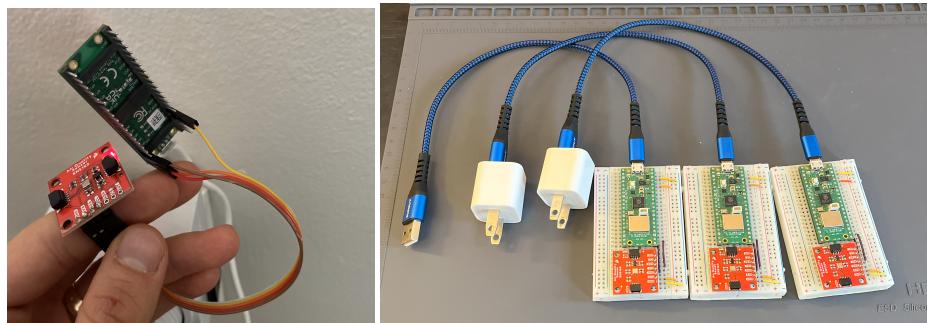


Figure 4: Development Unit (left) and Full Prototypes (right)

The software development for the sensor unit was relatively straight forward. We wrote simple programs to aid hardware bring up, first to blink the Pico LED

<sup>12</sup><https://github.com/EricSchrock/co2-monitor/blob/main/docs/co2-sensor.md>

<sup>13</sup><https://www.sciosense.com/wp-content/uploads/documents/SC-001224-DS-9-ENS160-Datasheet.pdf>

<sup>14</sup><https://electronics360.globalspec.com/article/17986/what-are-eco2-sensors>

<sup>15</sup><https://www.sparkfun.com/products/20844>

and then to turn the Pico LED on when we breathed on the CO2 sensor. Then we created a simple program that connects to wifi, connects to the CO2 sensor via I2c, connects to the server, and then periodically reads and reports the CO2 level (see the sensor flow chart above in the Technical Approach section).

### **Server**

todo: “This is where you give the details in your implementation. Talk about specific software packages you used, hardware modules, any algorithms or research papers you referred to, data structure and protocol choices, etc. You should provide at least an informal list of citations of all these external materials that went into your project.”

todo: Include picture of the Pi in a hub. Include a link to Pi website. Talk about how we chose the Pi (had it already through the class) (otherwise would have chosen cheaper model with less RAM). Talk about challenges with server process dying (solved with rc.local and cron job).

### **Website**

todo: “This is where you give the details in your implementation. Talk about specific software packages you used, hardware modules, any algorithms or research papers you referred to, data structure and protocol choices, etc. You should provide at least an informal list of citations of all these external materials that went into your project.”

todo: Talk through how we chose the web technology (Flask?). Include link to Flask website.

## **Results**

### **Project Timeline**

Overall, the project probably took a little over the projected 45 hours. The bookends, the initial design decisions and the project report, took much more time than planned, but the hardware bring up was exceptionally smooth and we did not need the planned second prototype. Our time estimates for the website were about right.

### **Project Objectives**

At a high level, we met the objective we laid out in our project proposal, which was “to create a network of CO2 sensors with a web interface” to allow “distributed monitoring of indoor air quality.” Our proposal also laid out the objectives of “simple, low cost, low energy” sensors.

We met the objective of simple sensors. We avoided extra functionality that would have increased cost and complexity (e.g. multi-function sensors or extra

website features). We also kept the setup instructions<sup>16</sup> simple (for a prototype).

As our design evolved from a battery powered sensor to a wall powered sensor, we dropped the objective of being low energy. This allowed us to make tradeoffs in favor of simplicity and speed. For example, it allowed us to keep our code simple by using wifi instead of BLE and by not implementing any low power modes. Now that we have a working prototype that has proved its value, power usage optimization would be a valuable consideration for a second prototype.

We also met the objective of low cost sensors, with each sensor unit costing roughly \$40. A productionized sensor unit, benefiting from economies of scale, could cost even less.

The base station prototype is pricier at \$88. This is largely because we used the 8 GB Raspberry Pi 4B because it is what we had on hand. Using the 1 GB model would drop the base station cost to \$48. A simple website doesn't take much compute power. A productionized version of the base station could use an even cheaper processor. It's even possible that one of the sensor units could function as the server for the rest of the sensor network.

### Sensor Unit

Part	Price*	Link
Raspberry Pi Pico WH	\$7.00	<a href="https://www.pishop.us/product/raspberry-pi-pico-wh/">https://www.pishop.us/product/raspberry-pi-pico-wh/</a>
12.5W power supply	\$8.00	<a href="https://www.pishop.us/product/raspberry-pi-12-5w-power-supply-us-white/">https://www.pishop.us/product/raspberry-pi-12-5w-power-supply-us-white/</a>
Sparkfun ENS160	\$19.95	<a href="https://www.sparkfun.com/products/20844">https://www.sparkfun.com/products/20844</a>
Half sized bread board	\$4.75	<a href="https://www.pishop.us/product/half-size-400-pin-diy-breadboard-white/">https://www.pishop.us/product/half-size-400-pin-diy-breadboard-white/</a>
Wires and headers	\$0.30	
<b>Total</b>	<b>\$40.00</b>	

\*Price does not include shipping or bulk discounts.

### Base Station

Part	Price*	Link
Raspberry Pi 4B (8 GB)	\$75.00	<a href="https://www.pishop.us/product/raspberry-pi-4-model-b-8gb/">https://www.pishop.us/product/raspberry-pi-4-model-b-8gb/</a>
15W power supply	\$8.00	<a href="https://www.pishop.us/product/raspberry-pi-15w-power-supply-us-white/">https://www.pishop.us/product/raspberry-pi-15w-power-supply-us-white/</a>

<sup>16</sup><https://github.com/EricSchrock/co2-monitor/blob/main/src/README.md>

Part	Price*	Link
Case	\$5.00	<a href="https://www.pishop.us/product/raspberry-pi-4-case-red-white/">https://www.pishop.us/product/raspberry-pi-4-case-red-white/</a>
<b>Total</b>	<b>\$88.00</b>	

\*Price does not include shipping or bulk discounts.

## CO2 Data Analysis

One question we had was whether our sensors would be accurate enough to provide useful data. We found that the data trends were reliable enough to interpret and act on. For example, below is around thirteen hours of CO2 data from three different rooms. Below that is a list tying trends in the data to my (Eric's) activity that day.

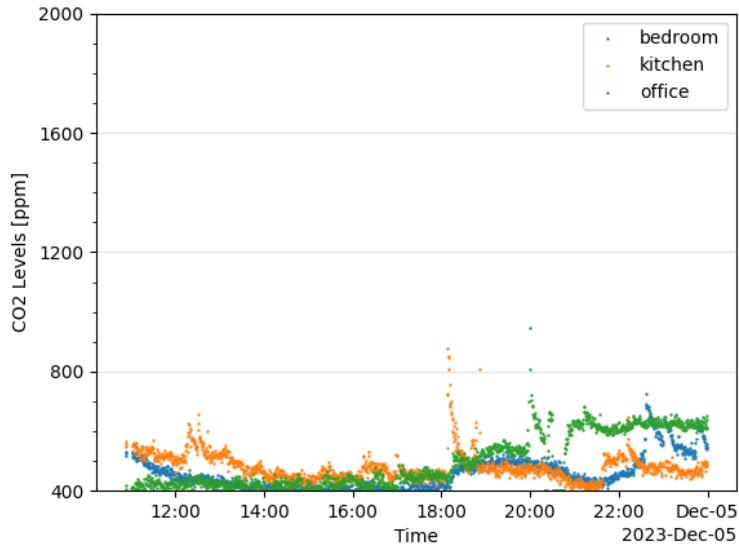


Figure 5: CO2 Data from Eric's Home (12/4/23)

1. 11 AM - 2 PM: The bedroom CO2 level drops from its nighttime high.
2. 12 - 1 PM: The kitchen CO2 level jumps as my wife and I eat lunch.
3. 1 - 6:15 PM: The bedroom sits empty and has the lowest CO2 level. I sit at my desk in my office and my wife moves around the kitchen and adjoining living room. The kitchen CO2 level is higher than the office CO2 level because humans breath out more CO2 when we are active.
4. 6:15 PM: The kitchen CO2 level spikes when we vent the pressure cooker and then dissipate to the rest of the house, leaving every room at a higher CO2 level.
5. 6:30 - 7 PM: The kitchen CO2 level jumps as my wife and I eat supper.
6. 7 - 9:30 PM: My wife leaves for time with friends and I return to my office. The kitchen CO2 level drops below that of the office.
7. 8:15 PM: I briefly open a window in my office. The office CO2 level drops to the minimum sensor value of 400 ppm, but jumps back up as soon as I close the window.
8. 8:30 - 8:45 PM: I open a window in my office for 15 minutes. The office

CO<sub>2</sub> level again drops to the minimum, but jumps back as soon as the window is closed.

9. 9:30 - 10:15 PM: My wife returns and we spend time together in the living room, near the kitchen. The kitchen CO<sub>2</sub> level jumps.
10. 10:15 PM: We go to bed. The kitchen CO<sub>2</sub> level drops and the bedroom CO<sub>2</sub> level jumps.

The sensors do have a couple issues to keep in mind. First, as explained in the section on implementation details, they are measuring eCO<sub>2</sub> which is an estimate of CO<sub>2</sub> levels derived from TVOC levels. TVOCs can spike without CO<sub>2</sub> spiking. For example, the data above shows a spike in the office CO<sub>2</sub> levels at around 8 PM. This was from running a paper shredder near the sensor.

Second, the sensor results can occasionally get stuck at a high offset until they are power cycled. The CO<sub>2</sub> level in the office jumped significantly the evening of December 4th (above) and remained high into the next day (below). I first tried opening windows throughout the house from 10:15 - 10:30 AM. The other rooms responded to this, but the office CO<sub>2</sub> level jumped back up as soon as the windows were closed. At around noon, I power cycled the office sensor.

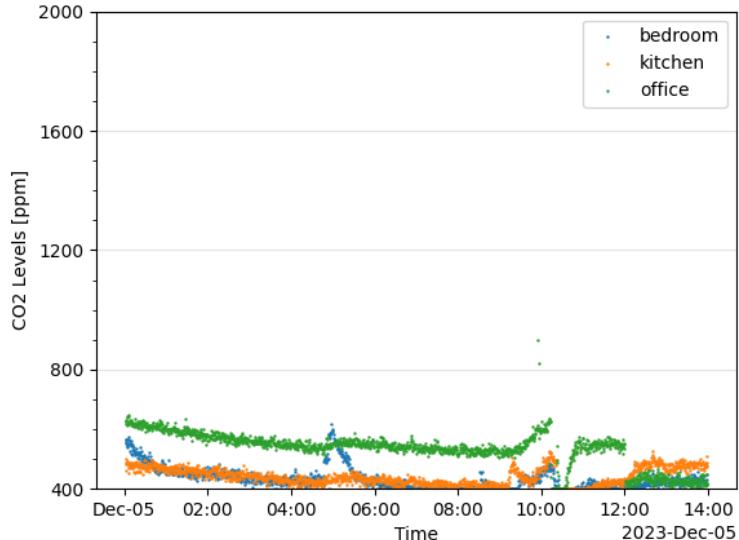


Figure 6: CO<sub>2</sub> Data from Eric's Home (12/5/23)

Despite these issues, we believe that the sensor data captured is accurate enough to drive decisions, such as whether to open windows or whether to invest in a new HVAC that cycles in outdoor air in response to high indoor CO<sub>2</sub> levels, as long as the sensor limitations are kept in mind.

todo (Devin): Do you have any interesting data or observations to add? How does your data compare to mine?

## Demo Videos

We recommend watching these videos on a large screen and/or setting your video player to HD/1080p.

- [Hardware Demo](#)
- [Website Demo](#)

todo (Eric): Record a demo video of how to setup the hardware.

- Sensor and base station hardware
- Installation (<https://github.com/EricSchrock/co2-monitor/blob/main/src/README.md>)
- Talk through LED transitions
- Show server log and CO2 data files

todo (Devin): Record a demo video of how to use the website.

- Installation (<https://github.com/EricSchrock/co2-monitor/blob/main/src/README.md>)
- Talk through and demo website features
  - Date selection
  - Sensor/room selection
  - Refresh for latest data
- Talk through interesting trends observed in the data

## Project Repository

<https://github.com/EricSchrock/co2-monitor>

## What We Learned

### Eric

Two of the key skills I learned in this project were networking via the Python sockets library and how to configure Linux using the `/etc/rc.local` file and `cron` jobs. Both were completely new to me and are valuable experiences for the future.

On a different note, I was surprised by how easy it was to get the project up and running on the Raspberry Pi Pico with MicroPython. I'm used to embedded boards being much more painful to bring up. Perhaps some personal projects I thought would be too time consuming are actually in reach!

This project also taught me about the health impacts of high concentrations of CO<sub>2</sub> in indoor spaces. This is important knowledge to help me stay sharp. As a remote worker, I can put what I've learned to use by opening my office windows at key times.

## **Devin**

I learned a great deal from working with the hardware in this lab. I soldered header pins for the first time. I gained valuable experience reading the ENS160 datasheet to understand how to read sensor values over I2C. I learned how to use `minicom` to view serial data from the Pi Pico and how to prototype micropython on the Pi Pico from the interactive micropython REPL.

This project significantly expanded my understanding and confidence in socket programming, which I had only just learned in previous class labs. It also challenged me to write a safe and robust server application.

Building the webpage may have taught me the most. Weighing different Python web frameworks gave great insight into the infrastructure underlying websites and web applications that I use on a daily basis. Actually deploying a functional web form involved more pieces than I knew and gave me a new respect for web development.

## **If We Had More Time**

If we were to take this project further, we would have three main goals. The first would be to increase the accuracy of the sensor readings, either by tuning them with temperature and humidity readings reported to them by the base station or by automatically power cycling the ENS160 CO<sub>2</sub> sensor periodically.

Second, we would expand the functionality of the website. We would add the ability to save timestamped notes to mark and explain phenomena in the data. We would also make the time interval for the display configurable. Additionally, we would add simple statistics, such as the max and average over the selected time interval.

The third goal would be to condense the sensor unit into a wall wart with a protective housing. We would look into 3D printing for the case and into a custom PCB to fit in a smaller form factor.

## **Conclusion**

In conclusion, we accomplished the high level goal of our project, to prototype a distributed, networked, and affordable CO<sub>2</sub> monitoring solution. Along the way, we tried new technologies, built new skills, and gained a deeper understanding of the quality of the air in our homes and of how it impacts our health.